

Christian Bernhard

christianbernhard089@gmail.com | +49 1778028796 | www.christian-bernhard.com

Summary

AI Engineer & Team Leader with 5+ years across software and AI, specialized in Generative AI, RAG, Vision-, Voice- and multi-agent systems — with a strong focus on scalable deployment and inference optimization on GPU-accelerated infrastructure. Hands-on with the NVIDIA AI stack (NIMs, TensorRT-LLM, vLLM) for production LLM workloads across cloud and on-prem, including hardware sizing, VRAM efficiency, and TCO. Proven track record designing and shipping enterprise-grade AI systems end-to-end. Strong in technical pre-sales and evangelism: 100+ workshops, talks, and trainings for enterprises, public institutions, and executives — equally comfortable with C-level stakeholders and engineering teams. Deep hands-on experience with AI-assisted development and coding agents (Codex, Claude Code, Cursor).

This CV captures the facts. To properly see my work, demos, and impact per role, visit christian-bernhard.com/about

Education

M. Sc. Computer Science - Ludwig-Maximilians-Universität München

2022 – 2024

- Thesis at BMW: Fine-Tuning Retriever and Generator Components in RAG: A Cross Domain-Study (GER: 1.4)

B. Sc. Business Informatics – University of Applied Sciences München

2018 - 2022

- Bachelor thesis: Natural Language Processing – Effort Estimation in Agile Projects (GER: 1.8)

Skills

Technologies: Python, Javascript, LangChain, LangGraph, HuggingFace, Transformers, PyTorch, FastAPI, Evaluation (LangSmith, Langfuse, DeepEval), IBM Watsonxai Orchestrate, Nvidia NIMs, Open & closed LLMs, VLM, Speech (GPT, LLaMA, Qwen, Nemotron, ...), Git, Azure, AzureAI, Docker, Kubernetes, CUDA & TensorRT Basics, Cursor, Claude Code, Codex

Work Experience

Teamlead AI – Cancom

05/2025 – now

- Leading AI initiatives and building CANCOM's enterprise AI portfolio, focusing on scalable LLM, RAG and agentic systems across industries
- Actively acting as a forward-deployed AI engineer and technical advisor, working closely with developers, architects and C-level stakeholders to design and deploy production-ready AI systems
- Acting as an AI enabler for enterprise customers, guiding end-to-end sovereign AI adoption – from initial edge deployments (e.g. DGX Spark) to fully scaled, data center-grade AI infrastructure (e.g. DGX B300 clusters)
- Hands-on deployment and optimization of LLM inference using NVIDIA AI Enterprise stack (NIMs, Nemotron, TensorRT-LLM, vLLM), including hardware sizing, GPU memory planning and performance tuning
- Building end-to-end AI applications (LLMs, Agents, VLMs, Speech2Speech) across cloud and on-prem environments, with focus on scalability and low-latency inference
- Creating technical demos, PoCs and sample applications (e.g. maintainer of: <https://demo.cancom-ds.de/>) to showcase AI capabilities and enable customer adoption
- Conducting AI Discovery Workshops, developer sessions and executive briefings to drive adoption of modern AI architectures and agentic workflows
- Collaborating with sales, engineering and partners to translate technical capabilities into business value and scalable solutions
- Mentoring AI engineers and contributing to internal enablement on topics such as GenAI, NVIDIA stack, and inference optimization

AI Researcher – BMW

10/2024 – 05/2025

- LoRA-based fine-tuning on small and large language models (e.g. LLaMA -3.1-8b, LLaMA-3.3.70b)
- Finetuning embedding models to enhance retrieval quality
- Evaluated end-to-end RAG systems using custom fine-tuned retrievers and LLMs
- Integrated LLM-as-a-Judge metrics and developed evaluation pipelines for automated quality scoring
- Leveraged state-of-the-art AI accelerator infrastructure (2 x NVIDIA H100 DGX Cluster) with TensorRT-LLM for efficient training and inference.

Head of Operations – Aqarios

04/2023 – 04/2025

- started as backend developer, developing different backend microservices
- led a team of 10+ software engineers and managed the development of the Luna Platform - a SaaS solution that enables businesses to seamlessly integrate with quantum computing backends and accelerates quantum application development: www.docs.aqarios.com
- lead our go to market strategy in 2024 and supervise our Marketing and HR department
- take part of conventions, represent the company, lead technical demos for our clients and partners like SAP, BASF, Siemens, BMW, etc.
- Conducted technical interviews and recruited software engineers

Founder – Ainleuchtend

09/2023 – now

- As part of my self-employment as an AI consultant, I advise companies on AI projects, develop individual AI software and hold training courses on AI throughout Germany.
- My clients include municipalities and public institutions, global companies such as Expedia, as well as private individuals – including former German President Christian Wulff. For more information: <https://www.ainleuchtend.de/>

Software Engineer – Adesso SE

09/2020 – 03/2023

- Developing a deep learning project; Developing and launching the iPadOS app “wryte” from scratch
- Team Leader for an international team, responsible for the App Store launch and updates
- appearance on the german tv-version of shark tank, temporary rank 1 Appstore in the education section, over 100k downloads

Engineering Fellow – Tech4Germany Fellowship Program (Berlin)

8/2022 – 11/2022

- 3 months fellowship program in Berlin - 32 talents develop prototypes for concrete digital projects of the federal administration. Developed a web-app dashboard for the Federal Ministry of Finance

Further Programming Expertise

Courses

I dedicate time daily to learning, primarily through hands-on projects in private repositories, thousands of hours of YouTube study, and my personal Obsidian knowledge base. Alongside this, I've completed 20+ AI certifications - on Coursera, DeepLearning.AI and NVIDIA's Deep Learning Institute - covering LLMs, RAG, Agents, Evaluation, LangChain, LangGraph, and MCP. Certificates: christian-bernhard.com/certificates

Furthermore

- Participated in the HackaTUM 2023 in Munich from the Technical University of Munich
- Participated in the Kaggle Challenge: Computer Vision Hacking the Human Vasculature
- LLM Implementation: Programmed LLaMA-2 from scratch in Python for research purposes
- Self-studying advanced GPU computing with a focus on GPU hardware fundamentals and NVIDIA AI Enterprise tools

Languages: German (Native), English (Fluent), Spanish (Beginner), Mandarin (Beginner)